

Machine Learning in Robotics

Lecture 2: Regression

Prof. Dongheui Lee

*Institute of Automatic Control Engineering
Technische Universität München*

dhlee@tum.de

Regression problems


- The goal is to make quantitative (real valued) predictions on the basis of a vector of features or attributes
- Examples: house prices, stock values, survival time, fuel efficiency of cars, etc.
- Questions: What can we assume about the problem? How do we formalize the regression problem? How do we evaluate predictions?

Linear Regression

- Linear regression assumes that expected value of the output given an input is linear.

$$y^{(i)} = wx^{(i)} + \epsilon \quad (1)$$

input x	output y
1	1
2	2
3	2.2
4	3.1
1.5	1.9




$x^{(i)}$: i-th input

$y^{(i)}$: i-th output

Linear Least Squares Regression : Single Parameter

- Which value of w makes the output values most likely?
- One that minimizes sum of squares of residuals.

$$E = \frac{1}{n} \sum_{i=1}^n (y^{(i)} - wx^{(i)})^2$$
$$w = \frac{\sum_{i=1}^n x^{(i)} y^{(i)}}{\sum_{i=1}^n x^{(i)2}}$$



- We can use it for prediction.

Linear Least Squares Regression

We need to define a class of functions (types of predictions we will try to make) such as linear predictions:

$$f(x) = w_0 + w_1x \quad (2)$$

where w_0 and w_1 are the parameters we need to set.

We need an estimation criterion so as to be able to select appropriate values for our parameters (w_0 and w_1) based on the training set $\{(x^{(i)}, y^{(i)}), \dots, (x^{(n)}, y^{(n)})\}$

For example, we can use the empirical loss:

$$E = \frac{1}{n} \sum_{i=1}^n (y^{(i)} - f(x^{(i)}))^2 \quad (3)$$

Estimating the parameters

- We minimize the empirical squared loss

$$\begin{aligned}
 E &= \frac{1}{n} \sum_{i=1}^n (y^{(i)} - f(x^{(i)}))^2 \\
 &= \frac{1}{n} \sum_{i=1}^n (y^{(i)} - w_0 - w_1 x^{(i)})^2
 \end{aligned}
 \tag{4}$$

- By setting the derivatives with respect to w_0 and w_1 to zero we get necessary conditions for the optimal parameter values

$$\begin{aligned}
 \frac{\partial}{\partial w_0} E &= 0 \\
 \frac{\partial}{\partial w_1} E &= 0
 \end{aligned}
 \tag{5}$$

Estimating the parameters

- By setting the derivatives with respect to w_0 and w_1 to zero

$$\begin{aligned}\frac{\partial}{\partial w_0} E &= 0 \\ \frac{\partial}{\partial w_1} E &= 0\end{aligned}\tag{6}$$

- we get necessary conditions for the optimal parameter values

$$\begin{aligned}w_0 &= \frac{n \sum x^{(i)} y^{(i)} - \sum x^{(i)} \sum y^{(i)}}{n \sum x^{(i)2} - (\sum x^{(i)})^2} \\ w_1 &= \frac{\sum y^{(i)} \sum x^{(i)2} - \sum x^{(i)} \sum x^{(i)} y^{(i)}}{n \sum x^{(i)2} - (\sum x^{(i)})^2}.\end{aligned}\tag{7}$$

Linear regression problem with multiple variables

We can express the solution a bit more generally by resorting to a matrix notation

$$X = \begin{bmatrix} \mathbf{x}^{(1)} \\ \mathbf{x}^{(2)} \\ \vdots \\ \mathbf{x}^{(n)} \end{bmatrix} = \begin{bmatrix} 1 & x_1^{(1)} & x_2^{(1)} & \cdots & x_m^{(1)} \\ 1 & x_1^{(2)} & x_2^{(2)} & \cdots & x_m^{(2)} \\ & & \vdots & & \\ 1 & x_1^{(n)} & x_2^{(n)} & \cdots & x_m^{(n)} \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(n)} \end{bmatrix} \quad \mathbf{w} = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_m \end{bmatrix}$$

so that $f(\mathbf{x}) = X\mathbf{w}$.

The result becomes

$$\mathbf{w}^* = (X^T X)^{-1} X^T \mathbf{y}$$

Solving Linear regression in matrix notation

Our empirical loss becomes $E = \frac{1}{n} \|y - Xw\|^2$.

By setting the derivatives of E with respect to w to zero, we get the same optimality conditions as before, now expressed in a matrix form

$$\begin{aligned}\frac{\partial}{\partial w} E &= \frac{1}{n} \frac{\partial}{\partial w} (y - Xw)^T (y - Xw) \\ &= \frac{1}{n} \frac{\partial}{\partial w} (w^T X^T X w - 2y^T X w + y^T y) \\ &= \frac{1}{n} \left(\frac{\partial w^T X^T X w}{\partial w} - 2y^T X \right) \\ &= \frac{1}{n} (2w^T X^T X - 2y^T X) = 0\end{aligned}$$

which yields

$$w^* = (X^T X)^{-1} X^T y$$

Gradient Descent

- Another way to minimize $E(\mathbf{w})$
- Start with an initial value of \mathbf{w} , keep changing \mathbf{w} to reduce $E(\mathbf{w})$

$$w_j := w_j - \alpha \frac{\partial}{\partial w_j} E(\mathbf{w})$$

$$w_j := w_j - 2\alpha(f(\mathbf{x}) - y)x_j$$

- Batch Gradient Descent
 - All training data is taken into account

$$w_j := w_j - \frac{\alpha}{n} \sum_{i=1}^n (f(\mathbf{x}^{(i)}) - y^{(i)})x_j^{(i)}$$

- Incremental (Stochastic) Gradient Descent

$$w_j := w_j - \alpha(f(\mathbf{x}^{(i)}) - y^{(i)})x_j^{(i)}, \text{ for } i = 1 \text{ to } n$$

Probabilistic approach

- Assume

$$y^{(i)} = \mathbf{x}^{(i)} \mathbf{w} + \epsilon^{(i)}$$

$$\epsilon^{(i)} \sim \mathcal{N}(0, \sigma^2)$$

$$p(y^{(i)} | \mathbf{x}^{(i)}; \mathbf{w}) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \mathbf{x}^{(i)} \mathbf{w})^2}{2\sigma^2}\right)$$

- Likelihood

$$\begin{aligned} L(\mathbf{w}) &= \prod_{i=1}^n p(y^{(i)} | \mathbf{x}^{(i)}; \mathbf{w}) \\ &= \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \mathbf{x}^{(i)} \mathbf{w})^2}{2\sigma^2}\right) \end{aligned}$$

- Choose parameters to maximize the likelihood
= same as minimizing LMS

Beyond linear regression

- The linear regression functions

$$f(\mathbf{x}) = w_0 + w_1x_1 + \dots + w_mx_m$$

are convenient because they are linear in the parameters, not necessarily in the input \mathbf{x}

- We can easily generalize these classes of functions to be non-linear functions of the inputs \mathbf{x} but still linear in the parameters \mathbf{w} . For example: m th order polynomial prediction

$$f(x) = w_0 + w_1x + w_2x^2 + \dots + w_mx^m$$

Quadratic Regression

$$f(\mathbf{x}) = w_0 + w_1x_1 + w_2x_2 + w_3x_1^2 + w_4x_1x_2 + w_5x_2^2$$

X_1	X_2	Y
3	2	7
1	1	3
\vdots	\vdots	\vdots

$\mathbf{X} =$	3	2	$\mathbf{y} =$	7
	1	1		3
	\vdots	\vdots		\vdots

$\mathbf{Z} =$	1	3	2	9	6	4	$\mathbf{y} =$	7
	1	1	1	1	1	1		3
	\vdots					\vdots		\vdots


$$\mathbf{z} = (1, x_1, x_2, x_1^2, x_1x_2, x_2^2)$$

$$\beta = (\mathbf{Z}^T \mathbf{Z})^{-1} (\mathbf{Z}^T \mathbf{y})$$

$$f(\mathbf{x}) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_1^2 + \beta_4 x_1 x_2 + \beta_5 x_2^2$$

Polynomial Curve Fitting





$$f(x) = w_0 + w_1x + w_2x^2 + \dots + w_mx^m$$



Minimize the empirical error


$$E = \frac{1}{n} \sum_{i=1}^n (y^{(i)} - f(x^{(i)}))^2$$

Polynomial Curve Fitting with different orders



Polynomial Curve Fitting



Root-mean-square error & Polynomial coefficients



	$m = 0$	$m = 1$	$m = 3$	$m = 9$
w_0^*	0.19	0.82	0.31	0.35
w_1^*		-1.27	7.99	232.37
w_2^*			-25.43	-5321.83
w_3^*			17.37	48568.31
w_4^*				-231639.30
w_5^*				640042.26
w_6^*				-1061800.52
w_7^*				1042400.18
w_8^*				-557682.99
w_9^*				125201.43

Polynomial Curve Fitting



9th order polynomials by increasing the training data, $n = 15$ and $n = 100$



Regularization


Penalize large coefficient values

$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^n (f(x^{(i)}, \mathbf{w}) - y^{(i)})^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$



Regularization

Root-mean-square error & Polynomial coefficients




	$\ln \lambda = -\infty$	$\ln \lambda = -18$	$\ln \lambda = 0$
w_0^*	0.35	0.35	0.13
w_1^*	232.37	4.74	-0.05
w_2^*	-5321.83	-0.77	-0.06
w_3^*	48568.31	-31.97	-0.05
w_4^*	-231639.30	-3.89	-0.03
w_5^*	640042.26	55.28	-0.02
w_6^*	-1061800.52	41.32	-0.01
w_7^*	1042400.18	-45.95	-0.00
w_8^*	-557682.99	-91.53	0.00
w_9^*	125201.43	72.68	0.01

Applet

<http://mste.illinois.edu/users/exner/java.f/leastquares/>

Phenological Application: Temperature-phenology relationship

Can we detect a response to temperature in phenology?




no.	Visualisation	Description	no.	Visualisation	Description
61		No flower petals visible	62		About 20 % of flowers open
68		Four flower petals visible (as published in Jochner)	63		About 30 % of flowers open
66		Four flowers open	64		About 60 % of flowers open
61		Beginning of flowering, about 50 % of flowers open or contrast to 60 distributed near the whole petal and more	65		Full Flowering or at least 50 % of flowers open, that petals falling



Jochner. 2008, Ellwood et al. 2013.

Robotics Applications: Odometry calibration

Estimate the pose (x, y, θ) of a mobile robot given the angular velocity of each wheel (ω_L, ω_R)



$$\dot{x} = v \cos(\theta)$$

$$\dot{y} = v \sin(\theta)$$

$$\dot{\theta} = \omega$$

$$\begin{bmatrix} v \\ \omega \end{bmatrix} = \mathbf{W}(r_R, r_L, b) \begin{bmatrix} \omega_R \\ \omega_L \end{bmatrix}$$

Odometry calibration consists in estimating \mathbf{W}

Robotics Applications: Odometry calibration

$$\begin{aligned}
 x_{t+1} &= x_t + v \cos(\theta) \Delta T \\
 y_{t+1} &= y_t + v \sin(\theta) \Delta T \\
 \theta_{t+1} &= \theta_t + \omega \Delta T
 \end{aligned}
 \Rightarrow
 \begin{aligned}
 \begin{bmatrix} x_N - x_0 \\ y_N - y_0 \end{bmatrix} &= \mathbf{X}(\omega_R, \omega_L) \begin{bmatrix} w_{11} \\ w_{12} \end{bmatrix} \\
 \theta_N - \theta_0 &= \mathbf{Z}(\omega_R, \omega_L) \begin{bmatrix} w_{21} \\ w_{22} \end{bmatrix}
 \end{aligned}$$

$$\begin{bmatrix} v \\ \omega \end{bmatrix} = \begin{bmatrix} w_{11} & w_{12} \\ w_{21} & w_{22} \end{bmatrix} \begin{bmatrix} \omega_R \\ \omega_L \end{bmatrix}$$

Executing M trajectories the parameters $(w_{11}, w_{12}, w_{21}, w_{22})$ can be identified using Linear Regression




Antonelli G., Chiaverini S., and Fusco G. *A Calibration Method for Odometry of Mobile Robots Based on the Least-Squares Technique: Theory and Experimental Validation*. Transactions on Robotics. 2005.

Robotics Applications: Obstacle avoidance

Assign an attractive potential (u_{att}) to the goal position and repulsive potentials (u_{rep}) to the obstacles


$$u = u_{att} + u_{rep} \rightarrow f_j = \frac{\partial}{\partial p_j} u_{att} + \frac{\partial}{\partial p_j} u_{rep}$$



Robotics Applications: Obstacle avoidance

A collision-free trajectory is generated using Gradient Descent

$$p^{(i+1)} = p^{(i)} - \alpha \frac{f^{(i)}}{\|f\|}$$



Khatib O. *Real-time obstacle avoidance for manipulators and mobile robots.*
International Journal of Robotics Research. 1986.



Reference

- Bishop, Pattern Recognition and Machine Learning
Chap. 1.1, 3, 6.4.1, 6.4.2
- Mitchell, Machine Learning
Chap. 4.4.3, 8.3