

Machine Learning in Robotics

Lecture 4: Unsupervised Clustering

Prof. Dongheui Lee

*Institute of Automatic Control Engineering
Technische Universität München*

dhlee@tum.de

Today Lecture Outline


- Unsupervised Clustering
- Similarity measures
- Criterion functions
- Iterative optimization algorithms
- K-means & Variations
- Hierarchical clustering

Supervised vs. Unsupervised learning


- Supervised learning
 - A pattern is a pair of variables $\{x, \omega\}$ where x is a collection of observations or features (feature vector) and ω is a the concept behind the observation (label)
- Unsupervised learning
 - Use unlabeled data, a collection of feature vectors without the class label ω
 - These methods are called unsupervised because they are not provided the correct answer
 - Unsupervised methods may appear to have limited capabilities, but they are useful
 - ▶ Labeling large data sets can be a costly procedure
 - ▶ Class labels may not be known beforehand
 - ▶ Large datasets can be compressed by finding a small set of proto-types

TEST : Unsupervised clustering

What is the natural grouping among those objects?



Many possibilities!!!



Two approaches for unsupervised learning

- Parametric approaches
 - Functional forms for the underlying class-conditional densities are assumed, and we must estimate the parameters

$$p(\mathbf{x}|\theta) = \sum_{i=1}^K p(\mathbf{x}|\omega_i, \theta_i) p(\omega_i)$$


- Non-parametric approaches
 - No assumptions are made about the underlying densities
 - Instead, seek a partition of the data into clusters
 - These methods are typically referred to as clustering

Vector Quantization (VQ)

- Vector Quantization is a lossy data compression method
- Mapping n feature vectors $X = \{x^{(1)}, x^{(2)}, \dots, x^{(n)}\}$ to K classes of feature vectors $Y = \{y^{(1)}, y^{(2)}, \dots, y^{(K)}\}$

$$y^{(j)} = c(x^{(i)}), \quad j = 1, \dots, K; \quad i = 1, \dots, n; \quad K < n$$

- Code vector $y^{(j)}$
 - Cluster centers
 - Points with high density in feature space
- Code book $Y = \{y^{(1)}, y^{(2)}, \dots, y^{(K)}\}$
 - The set of all code vectors



Nonparametric clustering

Nonparametric clustering involves three steps:

- Defining a measure of (dis)similarity between examples
- Defining a criterion (Distortion) function for clustering
- Defining an algorithm to minimize (or maximize) a criterion (distortion) function

Similarity Measures

- A measuring rule of $d(x, y)$ for the distance between two vectors x and y is considered a metric if it satisfies the following properties.
 - non-negativity: $d(x, y) \geq 0$
 - reflexivity: $d(x, y) = 0$ if and only if $x = y$
 - symmetry: $d(x, y) = d(y, x)$
 - triangle inequality: $d(x, y) + d(y, z) \geq d(x, z)$

Similarity Measures

- The most general form of distance metric is the power norm

$$d(\mathbf{x}, \mathbf{y}) = \left(\sum_{i=1}^m |x_i - y_i|^q \right)^{1/q}$$

- where $q \geq 1$ is a selectable parameter - general Minkowski metric
- When $q = 2$, Euclidean metric $d(\mathbf{x}, \mathbf{y}) = \|\mathbf{x} - \mathbf{y}\| = \sqrt{(\mathbf{x} - \mathbf{y})^T (\mathbf{x} - \mathbf{y})}$
- When $q = 1$, Manhattan or city block metric - sum of the absolute distances along each of the m coordinate axes.
- Notice that the above distance metrics are measures of dissimilarity.

Similarity Measures

- Inner product

$$s(\mathbf{x}, \mathbf{y}) = \frac{\mathbf{x}^T \mathbf{y}}{\|\mathbf{x}\| \|\mathbf{y}\|}$$

- When the features are binary-valued (0 or 1), the normalized inner product is a measure of the relative possession of common attributes.
- One variation is Tanimoto distance, a ratio of the number of shared attributes to the number possessed by \mathbf{x} or \mathbf{y} .

$$s(\mathbf{x}, \mathbf{y}) = \frac{\mathbf{x}^T \mathbf{y}}{\mathbf{x}^T \mathbf{x} + \mathbf{y}^T \mathbf{y} - \mathbf{x}^T \mathbf{y}}$$

Criterion (Distortion) function

- Once a (dis)similarity measure has been determined, we need to define a criterion function to be optimized.
- This function measures the clustering quality of any partition of the data.
 - The most widely used criterion function for clustering is the *sum-of-square-error*

$$J = \sum_{i=1}^K \sum_{x \in \omega_i} (x - y^{(i)})^2 \quad \text{where } y^{(i)} = \frac{1}{N_i} \sum_{x \in \omega_i} x$$

Here, N_i is the number of samples in the ω_i

- This criterion measures how well the data set is represented by the cluster centers
- Other criterion functions exist, based on the scatter matrices used in Linear Discriminant Analysis (LDA).
 - Trace criterion $tr[S_w]$
 - Determinant criterion $|S_w|$
 - Invariant criterion $tr[S_T^{-1} S_w]$

Iterative optimization

- Once a criterion function has been defined, we must find a partition of the data set that minimizes the criterion.
- Exhaustive enumeration of all partitions, which guarantees the optimal solution, is unfeasible.
- Common approach is to proceed in an iterative fashion
 - Find reasonable initial partition
 - Move samples from one cluster to another in order to reduce the criterion function
- These iterative methods produce sub-optimal solution but are computationally tractable
- Approaches
 - K-means algorithm
 - LBG algorithm
 - Fuzzy K-means algorithm
 - Basic iterative minimum-squared-error clustering

k-means, EM algorithm

- EM algorithm
 - Useful when estimating an optimal solution of a problem including hidden information
 - An iterative method which alternates between performing an expectation (E) step and a maximization (M) step
 - E step: computes the expectation of the log-likelihood evaluated using the current estimate for the latent variables
 - M step: computes parameters maximizing the expected log-likelihood found on the E step
 - Local convergence, No guarantee for global convergence
- K-means
 - A simple example of the EM optimization algorithm

k-means algorithm (Lloyd, 1982)

Cluster dataset $X = \{x^{(1)}, x^{(2)}, \dots, x^{(n)}\}$, given K

1. Initialization: Choose K random vectors $Y = \{y^{(1)}, y^{(2)}, \dots, y^{(K)}\}$ as an initial mean set
2. E-step: For each data point, find the closest class and label them. Dataset is divided into K classes

$$X_i = \left\{ x^{(j)} \mid d(x^{(j)}, y^{(i)}) \leq d(x^{(j)}, y^{(k)}), j = 1, \dots, n; k = 1, \dots, K \right\}$$

如果x距离y_i比距离其他的y的距离都近，则将这个x归类为i。共有K类。

3. M-step: From the current clusters, their mean vectors are updated


$$y^{(i)} = \frac{1}{N_i} \sum_{x \in \omega_i} x = C(X_i)$$

4. Calculate the total distortion, the sum of the distance between each datapoint and its closest cluster mean.


$$J = \sum_{i=1}^K \sum_{x \in \omega_i} (x - y^{(i)})^2$$

5. Evaluate the convergence. If converged, stop. Else, go to step 2.

Illustration of k-means algorithm



K=2



Blue : after E-step

Red : after M-step

Source: C. M. Bishop, Pattern Recognition and Machine Learning, Springer, 2006

K-means applet

Run K-means applet

http://home.deib.polimi.it/matteucc/Clustering/tutorial_html/AppletKM.html

Try 3 clusters, 100 points

k-means clustering applications



Source: C. M. Bishop, Pattern Recognition and Machine Learning, Springer, 2006

Basic Iterative Minimum-Squared-Error Clustering

- A sequential version of the k-means clustering algorithm.
 - k-means procedure waits until all n samples have been reclassified before updating.
 - the Basic Iterative Minimum-Squared-Error Clustering updates after each sample is reclassified.
- Disadvantages
 - more susceptible to being trapped in local minima
 - results depend on the order in which the candidates are selected
- Merits
 - at least a stepwise optimal procedure
 - suitable to problems in which samples are acquired sequentially and clustering must be done online

Basic Iterative Minimum-Squared-Error Clustering

Cluster dataset, given K

1. Initialization: Choose K random vectors $\mathbf{Y} = \{\mathbf{y}^{(1)}, \mathbf{y}^{(2)}, \dots, \mathbf{y}^{(K)}\}$ as an initial mean set
2. For each data point, find the closest class and label them.

$$i \leftarrow \arg \min_k d(\mathbf{x}^{(j)}, \mathbf{y}^{(k)})$$

3. If $N_i \neq 1$, then compute

$$\rho_j = \begin{cases} \frac{N_j}{N_j + 1} \|\mathbf{x} - \mathbf{y}^{(i)}\|^2 & \text{if } j \neq i \\ \frac{N_j}{N_j - 1} \|\mathbf{x} - \mathbf{y}^{(i)}\|^2 & \text{if } j = i \end{cases} \quad (1)$$

4. If $\rho_k \leq \rho_j$ for all j , then transfer \mathbf{x} to k -th cluster.
5. Recompute MSE, $\mathbf{y}^{(i)}, \mathbf{y}^{(k)}$.
6. Evaluate the convergence. If converged, stop. Else, go to step 2.

Fuzzy k-means algorithm

- Allows one piece of data to belong to two or more clusters
- Minimizing the below objective function

$$J = \sum_{k=1}^K \sum_{i=1}^n P(\omega_k | \mathbf{x}^{(i)})^m (\mathbf{x}^{(i)} - \mathbf{y}^{(k)})^2$$

- m is a parameter to adjust the blending of different clusters. If $m = 0$, k-means. For $m > 1$, the criterion allows belonging to multiple clusters
- $P(\omega_k | \mathbf{x}^{(i)})$ is the degree of membership of $\mathbf{x}^{(i)}$ in the cluster k
- an iterative optimization of the objective function shown above, with the update of membership $P(\omega_k | \mathbf{x}^{(i)})$ and the cluster centers $\mathbf{y}^{(k)}$ by:

$$P(\omega_k | \mathbf{x}^{(i)}) = \frac{1}{\sum_{j=1}^K \left(\frac{\|\mathbf{x}^{(i)} - \mathbf{y}^{(k)}\|}{\|\mathbf{x}^{(i)} - \mathbf{y}^{(j)}\|} \right)^{\frac{2}{m-1}}}$$
$$\mathbf{y}^{(k)} = \frac{\sum_{i=1}^n P(\omega_k | \mathbf{x}^{(i)})^m \mathbf{x}^{(i)}}{\sum_{i=1}^n P(\omega_k | \mathbf{x}^{(i)})^m}$$

Fuzzy k-means algorithm : Pseudocode

Cluster dataset, given K

1. Initialization: Initialize K random mean vectors $\mathbf{Y} = \{\mathbf{y}^{(1)}, \mathbf{y}^{(2)}, \dots, \mathbf{y}^{(K)}\}$. Initialize $P(\omega_k|\mathbf{x}^i)$ for $i = 1..n, k = 1..K$
2. Normalize $P(\omega_k|\mathbf{x}^{(i)})$ so that $\sum_k^K P(\omega_k|\mathbf{x}^{(i)}) = 1$ for $i = 1..n$
3. Update mean vectors

$$\mathbf{y}^{(k)} = \frac{\sum_{i=1}^n P(\omega_k|\mathbf{x}^{(i)})^m \mathbf{x}^{(i)}}{\sum_{i=1}^n P(\omega_k|\mathbf{x}^{(i)})^m}$$

4. Update the degree of membership of $\mathbf{x}^{(i)}$ in the cluster k

$$P(\omega_k|\mathbf{x}^{(i)}) = \frac{1}{\sum_{j=1}^K \left(\frac{\|\mathbf{x}^{(i)} - \mathbf{y}^{(k)}\|}{\|\mathbf{x}^{(i)} - \mathbf{y}^{(j)}\|} \right)^{\frac{2}{m-1}}}$$

5. If $\max_{i,k} P(\omega_k|\mathbf{x}^{(i)})$ is converged, then stop. Else, go to step 2.

Fuzzy K-means applet

Run Fuzzy K-means applet


http://home.deib.polimi.it/matteucc/Clustering/tutorial_html/AppletFCM.html

Another variation of k-means

- ISODATA, which stands for Iterative Self-Organizing Data Analysis Technique (Algorithm) is an extension to the k-means algorithm with some heuristics to automatically select the number of clusters
- The algorithm works in an iterative fashion
 - (1) Perform k-means clustering
 - (2) Split any clusters whose samples are sufficiently dissimilar
 - (3) Merge any two clusters sufficiently close
 - (4) Go to (1)


Hierarchical clustering

- k-means and ISODATA create disjoint clusters, resulting in a "flat" data representation
 - ▶ However, sometimes it is desirable to obtain a hierarchical representation of data, with clusters and sub-clusters arranged in a tree-structured fashion (i.e., biological taxonomy)
- Hierarchical clustering methods can be grouped in two general classes
 - ▶ Agglomerative (bottom-up, merging) : Starting with n singleton clusters, successively merge clusters until one cluster is left
 - ▶ Divisive (top-down, splitting) : Starting with a unique cluster, successively split the clusters until n singleton examples are left



Dendrograms

- The preferred representation for hierarchical clusters is the dendrogram
- The dendrogram is a binary tree that shows the structure of the clusters
 - ▶ In addition to the binary tree, the dendrogram provides the similarity measure between clusters (the vertical axis)
- An alternative representation is based on sets
 - ▶ $\{\{x_1, \{x_2, x_3\}\}, \{\{\{x_4, x_5\}, \{x_6, x_7\}\}, x_8\}\}$
 - ▶ However, unlike the dendrogram, sets cannot express quantitative information



Divisive clustering

1. Start with one large cluster
 2. Find "worst" cluster
 3. Split it
 4. If $K < n$, go to step 2
- How to choose the "worst" cluster
 - How to split clusters
 - The computations required by divisive clustering are more intensive than for agglomerative clustering methods.

Non-uniform Binary Split Algorithm

1. Initialization: Calculate a center for all data points. $k := 1$ k is the class number

$$\mathbf{y}^{(1)} = \mathcal{C}(\mathbf{X})$$

2. Choose a class which has the largest distortion among current classes

$$J_i \geq J_j, \quad \forall j = 1, \dots, k$$

3. Split the class into two subclasses by using a small random vector

$$\mathbf{X}_a = \left\{ \mathbf{x} \text{ which is closer to } \mathbf{y}^{(i)} + \mathbf{v}_i \right\}$$


$$\mathbf{X}_b = \left\{ \mathbf{x} \text{ which is closer to } \mathbf{y}^{(i)} - \mathbf{v}_i \right\}$$

4. Update the code vectors


$$\mathbf{y}^{(i)} := \mathcal{C}(\mathbf{X}_a) \quad \text{and} \quad \mathbf{y}^{(k+1)} := \mathcal{C}(\mathbf{X}_b)$$

5. If $k = K$, stop. Else, $k := k + 1$ and go to step 2.

k-means vs. Non-uniform binary split



(a) k-means



(b) Binary split

Non-uniform binary split: Not good quality of codebook. But very fast

LBG algorithm

- Combination of k-means and Binary split
- Proposed by Linde, Buzo, and Gray (1980)
- Instead of random initial codebook, let's use codebook from binary split
- Faster convergence and better quality codebook than k-means algorithm

Pseudo-code of LBG algorithm

1. Initialization: Calculate a center for all data points. $k := 1$

$$\mathbf{y}^{(1)} = \mathcal{C}(\mathbf{X})$$

2. Splitting: For each cluster, split it into two subclasses. $k := 2k$


$$\mathbf{y}_a^{(i)} = \mathbf{y}^{(i)} + \mathbf{v}_i \text{ and } \mathbf{y}_b^{(i)} = \mathbf{y}^{(i)} - \mathbf{v}_i$$

3. Iteration (k-means algorithm)


- a E-step: For each data point, find the closest cluster which achieves the minimum distance measure
- b M-step: From the current cluster, update their mean vectors
$$\mathbf{y}^{(i)} = \mathcal{C}(\mathbf{X}_i)$$
- c Iterate E-step and M-step until it converges

4. If the desired number of code vectors is obtained ($k = K$), stop. Else, go to step 2.


Comparison



kmeans



Binary Split



LBG

Agglomerative clustering

1. Start with n singleton cluster
 2. Find nearest clusters
 3. Merge them
 4. If $K > 1$, go to step 2
- How to find the "nearest" pair of clusters
 - ▶ Minimum distance $d_{min}(X_i, X_j) = \min_{\mathbf{x} \in \omega_i, \mathbf{x}' \in \omega_j} \|\mathbf{x} - \mathbf{x}'\|$
 - ▶ Maximum distance $d_{max}(X_i, X_j) = \max_{\mathbf{x} \in \omega_i, \mathbf{x}' \in \omega_j} \|\mathbf{x} - \mathbf{x}'\|$
 - ▶ Average distance $d_{avg}(X_i, X_j) = \frac{1}{N_i N_j} \sum_{\mathbf{x} \in \omega_i} \sum_{\mathbf{x}' \in \omega_j} \|\mathbf{x} - \mathbf{x}'\|$
 - ▶ Mean distance $d_{mean}(X_i, X_j) = \|\mathbf{y}^{(i)} - \mathbf{y}^{(j)}\|$

Agglomerative clustering

Minimum distance

- When d_{min} is used to measure distance between clusters, the algorithm is called the *nearest neighbor* or *single-linkage* clustering algorithm
- If the algorithm is allowed to run until only one cluster remains, the result is a minimum spanning tree (MST)

Maximum distance

- When d_{max} is used to measure distance between clusters, the algorithm is called the *farthest neighbor* or *complete-linkage* clustering algorithm
- From a graph-theoretic point of view, each cluster constitutes a complete sub-graph


Average and mean distance

- The average and mean distance approaches are more robust to outliers
- Of the two, the mean distance is computationally more attractive, since Notice that the average distance approach involves the computation of $N_i N_j$ distances for each pair of clusters

Agglomerative clustering example

Perform agglomerative clustering on the following dataset using the single-linkage metric

- $X = \{1, 3, 4, 9, 10, 13, 21, 23, 28, 29\}$
- In case of ties, always merge the pair of clusters with the largest mean
- Indicate the order in which the merging operations occur



Agglomerative clustering applet

Run Agglomerative clustering applet



- single-linkage clustering
- complete-linkage clustering
- Average-linkage clustering

http://home.deib.polimi.it/matteucc/Clustering/tutorial_html/AppletH.html

Multi-Robot Exploration

- Task definition
 - ▶ Exploration of unknown areas by means of k mobile robots.
 - ▶ Application: Search/Rescue robot, planetary exploration, reconnaissance
- Problems
 - ▶ Reduce the difference of waiting time among different regions of a workspace.
 - ▶ Ensure a balanced exploration of the environment

Multi-Robot Exploration



- K-means is used to partition an unexplored space in regions.
- Each region is assigned to a robot solving an optimization problem.




Wu L., Puig D., and Garcia M. A. *Balanced Multi-Robot Exploration through a Global Optimization Strategy*. Journal of Physical Agents. 2010.



Outdoor Robots Navigation


- Task definition
 - ▶ Autonomous navigation for outdoor, unstructured environments.
 - ▶ Recognize navigable terrain and avoid obstacles, based on their appearance.
- Problem
 - ▶ Do reliable segmentation of outdoor scenes in an efficient manner (online).

Outdoor Robots Navigation




- Use compact texture/color descriptors (feature vectors) and fast unsupervised clustering algorithms.
- K-means is used to cluster neighborhood features vectors in a small set of basis vectors (*textons*).
- Each pixel is classified as belonging to one cluster using Euclidean distance (b).

Outdoor Robots Navigation



- Extract an histogram counting textons in a neighborhood of each pixel
- K-means to extract a set of histogram profiles ($k = 8$).
- *Earth Movers Distance* is used to merge similar clusters (c).

 Blas M. R., Agrawal M., Sundaresan A., and Konolige K. *Fast Color/Texture Segmentation For Outdoor Robots*. IEEE International Conference on Intelligent Robots and Systems. 2008.

Summary and Next Lecture

- What we have learned
 - k-means algorithm
 - hierarchical clustering
- Reading: Duda Chap. 10.4, 10.6-9, Bishop Chap. 9.1, Michell Chap. 6.12
- Next Lecture
 - Maximum Likelihood Estimation
 - Gaussian Mixture Model