

Machine Learning in Robotics

Lecture 11: Introduction to Reinforcement Learning

Prof. Dongheui Lee

*Institute of Automatic Control Engineering
Technische Universität München*

dhlee@tum.de


Reinforcement Learning

Learning of a behavior without explicit information about correct actions

- Between supervised and unsupervised learning
- No training patterns, but **rewards**
- Inspired by principles of human and animal learning
- Mild assumptions on the process to be controlled
- A control strategy can be learned from scratch

Architecture

The agent-environment interaction in reinforcement learning



The Environment

- The environment contains the process to be controlled
- Markov Decision Process (MDP): The environment is modeled by an MDP which is tuple $(S, A, \{P_{sa}\}, \gamma, R)$
 - ▶ S is a set of **states**
 - ▶ A is a set of **actions**
 - ▶ P_{sa} are the **state transition probabilities**.
 - ▶ $\gamma \in [0, 1)$ is the **discount factor**.
 - ▶ $R : S \times A \mapsto \mathbb{R}$ is the **reward function** (Rewards can also be a function of state S only and in that case $R : S \mapsto \mathbb{R}$).

Task for the Agent

Find a behavior which maximizes the expected total reward

For how long should we consider?

Finite Horizon

$$\max \left[\sum_{t=0}^T r_t \right]$$

Infinite Horizon

$$\max \left[\sum_{t=0}^{\infty} \gamma^t r_t \right]$$

γ is a discount factor ($0 \leq \gamma < 1$)

Reward function

The reward function controls which task should be solved

- Game (Checkers, chess)
Reward only at end: +1 when winning, -1 when loosing
- Avoiding mistakes (pole balancing)
Reward -1 at the end (when falling)
- Find a fast/short/cheap path to a goal
Reward -1 at each step

Simplifying assumptions

- Discrete time
- Finite number of actions $a_i \in a_1, a_2, a_3, \dots, a_n$
- Finite number of states $s_i \in s_1, s_2, s_3, \dots, s_m$
- Environment is a stationary markov decision process
- Reward r only depends on s

Policy and Value function

- Policy

Policy provides a mapping from states to action.

$$\pi(s) \mapsto a$$

- Value Function

Expected total future reward when starting from s and following policy π

$$\begin{aligned} V^\pi(s) &= E[R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots | s_0 = s, \pi] \\ &= R(s) + \gamma \sum_{s' \in S} P_{s\pi(s)}(s') V^\pi(s') \quad (\text{Bellman's equation}) \end{aligned}$$

Optimal Policy

An optimal policy is the the one which maximizes the value function

$$V^*(s) = R(s) + \max_{a \in A} \gamma \sum_{s' \in S} P_{sa}(s') V^*(s')$$

$$\pi^*(s) = \arg \max_{a \in A} \gamma \sum_{s' \in S} P_{sa}(s') V^*(s')$$

Classical problem: Grid world

- 11 states. Each **state** is represented by a position in the grid world.
- The agent **acts** deterministically by moving to other position.
 $A = \{N, S, E, W\}$
- reward: $R(4, 3) = 1$, $R(4, 2) = -1$, $R(s) = -0.02$ for all other states
- transition probability: 0.8 for a planned state and 0.1 for the other adjacent two states.

Value Iteration

For each state s , initialize $V(s) := 0$.

Repeat until convergence

{

For every state, update $V(s) := R(s) + \max_{a \in A} \gamma \sum_{s'} P_{sa}(s') V(s')$

}

$V(s)$ can be updated in synchronous and asynchronous manner.

Policy Iteration

Initialize π randomly.
Repeat until convergence
{
(a) Let $V := V^\pi$
(b) For each state s , let $\pi(s) := \arg \max_{a \in A} \sum_{s'} P_{sa}(s') V(s')$
}

Step (a) can be calculated by solving linear equations (with equal number of equations and unknowns).

Monte-Carlo Method

Start at some random state.

Follow π , store the rewards and s_t .

When the goal is reached, update $V^\pi(s)$ estimation for all visited states with the future reward we actually received.

- Monte-Carlo method is suitable only for episodic tasks
- Learns incrementally from episode-by-episode but not step-by-step

Temporal Difference Learning

There are two estimates of the value of a state:

- Before: $V^\pi(s_t)$
- After: $R_{t+1} + \gamma V^\pi(s_{t+1})$


Temporal Difference Learning

Idea: The second estimate is better!

$$V^{\pi}(s_t) \leftarrow V^{\pi}(s_t) + \alpha(R_{t+1} + \gamma V^{\pi}(s_{t+1}) - V^{\pi}(s_t))$$

- Learns considerably faster than the Monte-Carlo method
- Step by step learning.

Eligibility Trace



Q-Learning

Whenever reward r or next state s' cannot be predicted, we cannot calculate π even with a good estimate for V

$Q^\pi(s, a)$, is the expected infinite-horizon discounted return for executing a in state s and thereafter following π

$$Q^\pi(s, a) = E [R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | s_t = s, a_t = a, \pi]$$

$$\pi(s) = \arg \max_a Q(s, a)$$

$$V^*(s) = \max_a Q^*(s, a)$$

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [R_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t)]$$


Autonomous helicopter flight via RL




Value of each policy is calculated through Monte-Carlo method *PEGASUS*

method uses the observation that almost all computer simulations sample $s' \sim P_{sa}(\cdot)$ by first calling a random number generator to get one (or more) random numbers p , and then calculating s' as some deterministic function of the input s, a and the random p Since the helicopters model is stochastic,

random number were fixed in advanced to evaluate different policies

 HJ Kim, Michael I Jordan, Shankar Sastry, and Andrew Y Ng., *Autonomous helicopter flight via reinforcement learning*, In Advances in neural information processing systems, 2003.

Learning Motor Primitives using Reinforcement Learning



POWER is an Expectation Maximization based RL algorithm which does not require learning rate as a parameter:

$$\theta' = \theta + \frac{E\{\sum_{t=1}^T \varepsilon_t Q^\pi(s_t, a_t, t)\}}{E\{\sum_{t=1}^T Q^\pi(s_t, a_t, t)\}} \quad \text{where } \varepsilon_t \text{ is exploration term}$$



Jens Kober and Jan Peters, *Learning motor primitives for robotics*, pp. 2112 - 2118, ICRA, 2009.



Reading Material

- Mitchell, Chapter 13
- Russell and Norvig, Artificial Intelligence: A Modern Approach, Chapter 21
- Sutton and Barto, Reinforcement Learning: An Introduction